

# Document In-Memory Persistent Database

Mahdy Mousa Hamad

# Why Build a Database Management System?

- Educational Purposes (Research)
- Paid Service (Business Value)
- Doing Something New
- OpenSource




# Document Database Requirements

- Authorization and Authentication
- Flexibility in storing records
- Efficient Queries
- Horizontal Scalability
- Data Sharding
- Persistent Data on-disk
- ... and more....

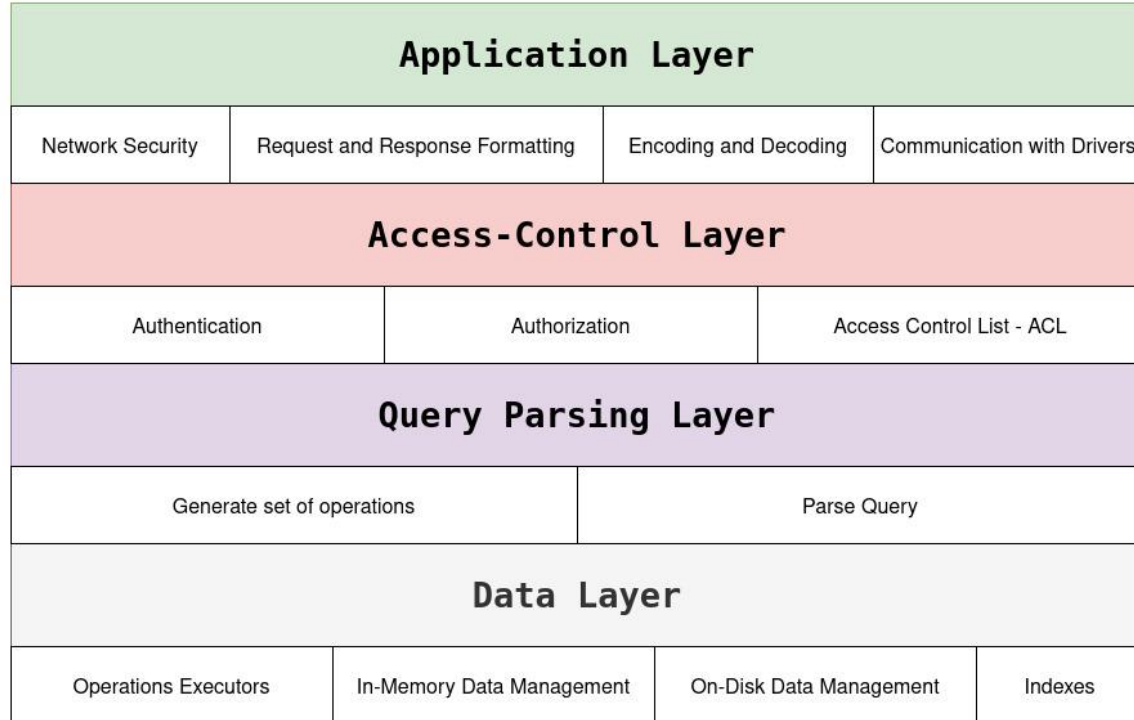


# Project Scope

- Authentication and Authorization - ACL
- Efficient Queries
- Persistent Data on-disk
- In-Memory Data Management and Indexing
- Simple On-Disk Indexing
- Exposed Network Layer for outside communication
- Simple Shell Client



# Database Architecture - Overview





# How Data is Queried Efficiently?


This is done using ***indexing***. In database systems, indexing is a way of storing data (usually in a tree structure) in a way that increase the performance of retrieving data that is associated with a key.

Indexes can be implemented on-disk and in-memory. We will focus on indexes stored in-memory. So data will be loaded into memory as indexes using BSTs (Binary Search Trees).

# Binary Search Tree - Indexing



# In-Memory Indexing






# On-Disk Indexing Simple Indexing

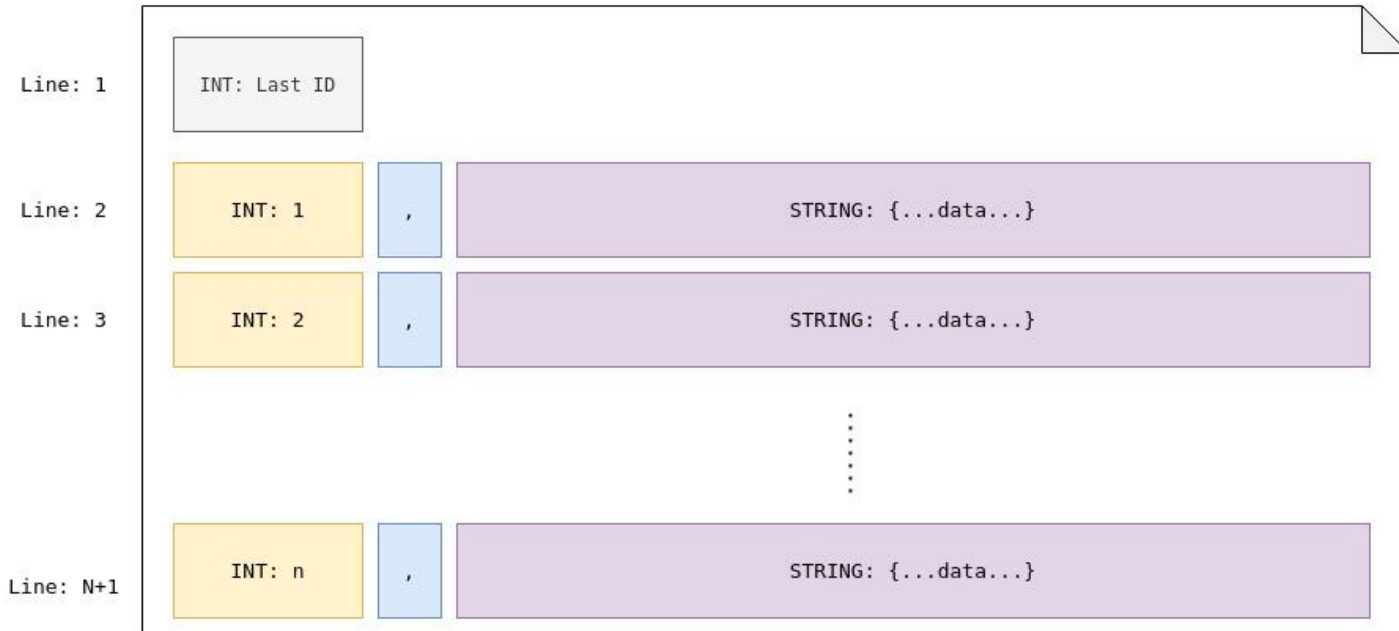
Find 2

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 |
| 1 | 2 | 3 | 4 | 5 |



# On-Disk Indexing Simple Indexing

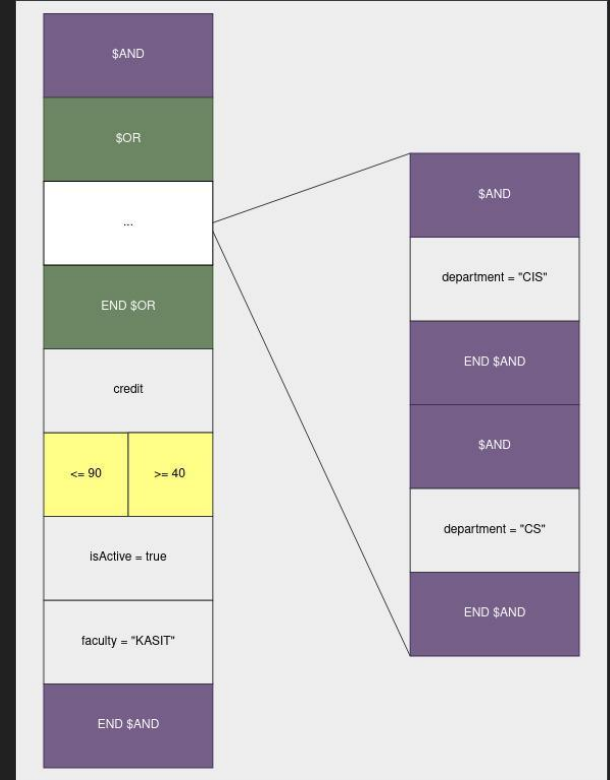
## Collection File Structure



# Parsing Query - Operations Builder

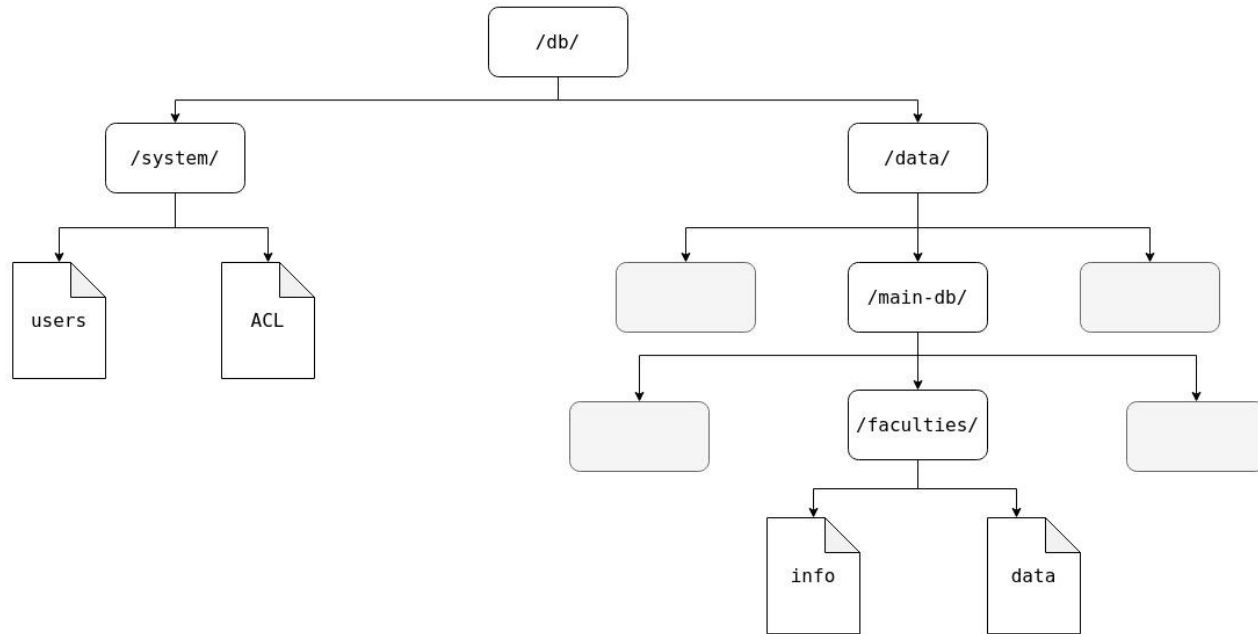
```
{  
  "faculty": "KASIT",  
  "isActive": true,  
  "$OR": [  
    {"department": "CS"},  
    {"department": "CIS"}  
  ],  
  "credit": {  
    "$lte": 90,  
    "$gte": 40  
  }  
}
```

Raw JSON Query



Operations Stack

# Database Storage - File Structure



# Implementation - Tools and Technologies



Database implementation using pure Java

- + Disk Management
- + In-Memory Indexing
- + Shell Client
- + Multi-Threading Programming
- + Socket Programming - Networking



Source Code Hosting

- + Version Control
- + Open-source the database

## Students Data with + 1 Million Records

## Students Data with + 1 Million Records

# Sample Document

```
{  
  "creditHours":1,  
  "firstName":"FirstName1",  
  "lastName":"LastName1",  
  "gender":"f",  
  "year":2001,  
  "gpa":1,  
  "_id":1,  
  "department":"bit",  
  "username":"test1",  
  "faculty":"KASIT",  
  "status":"I"  
}
```

# Let Us Query Some Data

- 1- Get all students with ***GPA greater than 3 and in the CS or in CIS department.***
- 2- Get all students that ***joined on 2010 and gpa equals 4*** (without index)
- 3- Get all students that ***joined on 2010 and gpa equals 4*** (using index)



# Let Us Query Some Data

Get all students with ***GPA greater than 3*** and ***in the CS or CIS department***  
(using 2 indexes: **department index** and **gpa index**)

```
maindb.students.find({"$OR":[{"department":"cs"}, {"department":"cis"}], "gpa":{"$gte":3}})
Query: {"$OR":[{"department":"cs"}, {"department":"cis"}], "gpa":{"$gte":3}}
Query from index
Operations: [department=cis]
Count: 75000
Query from index
Operations: [department=cs]
Count: 125000
Query from index
Operations: [gpa>=3]
Count: 200000
Total Count: 25000
<QuerySet [{"creditHours":8,"firstName":"FirstName8","lastName":"LastName8","gender":"m","year":2008,
...(remaining elements truncated)...>
Jan 17, 2023 6:17:32 PM dataLayer.Resolver resolve
INFO: SELECT query resolved in 98 milli seconds
```

# Let Us Query Some Data

Get all students that *joined on 2010* and *gpa equals 4* (without index)

```
maindb.students.find({"year":2010, "gpa":4})
```

```
Query: {"year":2010, "gpa":4}
```

```
Full collection search
```

```
Operations: [gpa=4, year=2010]
```

```
Count: 8695
```

```
<QuerySet [{"creditHours":79,"firstName":"FirstName79","last  
...(remaining elements truncated)...}>
```

```
Jan 17, 2023 5:24:54 PM dataLayer.Resolver resolve
```

```
INFO: SELECT query resolved in 7624 milli seconds
```

# Let Us Query Some Data

Get all students that *joined on 2010* and *gpa equals 4* (using index)

```
maindb.students.find({"year":2010, "gpa":4})
```

```
Query: {"year":2010, "gpa":4}
```

```
Query from index
```

```
Operations: [gpa=4, year=2010]
```

```
Count: 8695
```

```
Jan 17, 2023 5:21:35 PM dataLayer.Resolver resolve
```

```
INFO: SELECT query resolved in 7 milli seconds
```

```
<QuerySet [{"creditHours":79,"firstName":"FirstName79","last  
...(remaining elements truncated)...}>
```

# Future Work and Enhancements

This database is far from perfect, it needs some fundamental enhancements:

- Handling Concurrency Connections
- Horizontal Scaling
- Data Sharding
- On-Disk Indexing
- Building Native Drivers for Programming Languages

Questions?

Thank You